

# Diagnostic test evaluation with imperfect reference test

---

Paolo Eusebi

16/09/2021

# Recap

- Fitting models using MCMC is easy with JAGS / runjags
- But we must **never forget** to check convergence and effective sample size!
- More complex models become easy to implement
- So how do we extend these models to multiple diagnostic tests?

## **Hui-Walter models for 2 tests and 1 population**

---

# Hui-Walter Model

- A particular model formulation that was originally designed for evaluating diagnostic tests in the absence of a gold standard
- Not necessarily (or originally) Bayesian but often implemented using Bayesian MCMC
- But evaluating an imperfect test against another imperfect test is a bit like pulling a rabbit out of a hat
- If we don't know the true disease status, how can we estimate sensitivity or specificity for either test?

## Hui-Walter Model: 2 tests, 1 population

- D+

	T2+	T2-
T1+	$p * Se1 * Se2$	$p * Se1 * (1 - Se2)$
T1-	$p * (1 - Se1) * Se2$	$p * (1 - Se1) * (1 - Se2)$

- D-

	T2+	T2-
T1+	$p * Sp1 * Sp2$	$p * Sp1 * (1 - Sp2)$
T1-	$p * (1 - Sp1) * Sp2$	$p * (1 - Sp1) * (1 - Sp2)$

- 5 parameters: Se1, Se2, Sp1, Sp2, p

## Hui-Walter Model: 2 tests, 1 population

- We don't know the disease status » Marginalize over  $D+$  and  $D-$

	T2+	T2-
T1+	$p * Se1 * Se2 + p * Sp1 * Sp2$	$p * Se1 * (1 - Se2) + p * Sp1 * (1 - Sp2)$
T1-	$p * (1 - Se1) * Se2 + p * (1 - Sp1) * Sp2$	$p * (1 - Se1) * (1 - Se2) + p * (1 - Sp1) * (1 - Sp2)$

- 5 parameters and 3 equations » Non-identifiable model!
- We can't fit this model with standard likelihood procedures.
- Let's try to go Bayesian!

## Model Specification

```
model{  
  y ~ dmulti(prob, TotalTests)  
  
  # Test1- Test2-  
  prob[1] <- (p*((1-se[1])*(1-se[2]))) + ((1-p)*((sp[1])*(1-sp[2])))  
  # Test1+ Test2-  
  prob[2] <- (p*((se[1])*(1-se[2]))) + ((1-p)*((1-sp[1])*(1-sp[2])))  
  # Test1- Test2+  
  prob[3] <- (p*((1-se[1])*(se[2]))) + ((1-p)*((sp[1])*(se[2])))  
  # Test1+ Test2+  
  prob[4] <- (p*((se[1])*(se[2]))) + ((1-p)*((1-sp[1])*(se[2])))
```

## Model Specification

```
p ~ dbeta(1, 1)
se[1] ~ dbeta(1, 1)
sp[1] ~ dbeta(1, 1)
se[2] ~ dbeta(1, 1)
sp[2] ~ dbeta(1, 1)
#data# y, TotalTests
#monitor# p, prob, se, sp
#inits# p, se, sp
}
```



## Data, initial values, run

	T2+	T2-
T1+	48	4
T1-	12	36

```
y <- c(48, 12, 4, 36)
TotalTests <- sum(y)
p <- list(chain1=0.05, chain2=0.95)
se <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
sp <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
results <- run.jags(hw_2t_1p, n.chains=2)
```

[Remember to check convergence and effective sample size!]

## results

	Lower95	Median	Upper95	SSeff	psrf
p	0.327	0.499	0.667	4240	2.240
prob[1]	0.366	0.461	0.557	13858	1.000
prob[2]	0.074	0.133	0.202	13923	1.001
prob[3]	0.018	0.055	0.103	10259	1.000
prob[4]	0.257	0.344	0.439	13886	1.000
se[1]	0.026	0.548	1.000	4561	15.217
se[2]	0.000	0.389	0.966	4497	13.519
sp[1]	0.000	0.469	0.973	4609	15.081
sp[2]	0.037	0.620	1.000	5016	13.454

- Note the wide credible intervals!

- These models need a lot of data and/or strong priors for one of the tests
- Convergence is more problematic than usual
- Be very careful with the order of combinations in `dmultinom!`
- Check your results carefully to ensure they make sense!

## Label Switching

How to interpret a test with  $Se=0\%$  and  $Sp=0\%$ ?

The test is perfect - we are just holding it upside down...

We can force  $Se + Sp \geq 1$ :

```
se[1] ~ dbeta(1, 1)
sp[1] ~ dbeta(1, 1)T(1-se[1], )
```

or:

```
se[1] ~ dbeta(1, 1)T(1-sp[1], )
sp[1] ~ dbeta(1, 1)
```

But not both!

This allows the test to be useless, but not worse than useless

## Simulating data

Analysing simulated data is useful to check that we can recover parameter values.

```
se1 <- 0.9; sp1 <- 0.95
se2 <- 0.8; sp2 <- 0.99
prevalence <- 0.5; N <- 100
truestatus <- rbinom(N, 1, prevalence)
Test1 <- rbinom(N, 1, (truestatus * se1) +
                    ((1-truestatus) * (1-sp1)))
Test2 <- rbinom(N, 1, (truestatus * se2) +
                    ((1-truestatus) * (1-sp2)))
t <- table(Test1, Test2)
y <- as.numeric(t)
```

Can we recover these parameter values?

## Exercise

- Modify the code in the Hui-Walter model to force tests to be no worse than useless
- Simulate data and recover parameters for  $N=10$ , 100 or 10000

## **Hui-Walter models for 2 tests and N populations**

---

## Hui-Walter : 2 tests, 2 pops

- Population 1

	T2+	T2-
T1+	$p1*Se1*Se2+p1*Sp1*Sp2$	$p1*Se1*(1-Se2)+p1*Sp1*(1-Sp2)$
T1-	$p1*(1-Se1)*Se2+p1*(1-Sp1)*Sp2$	$p1*(1-Se1)*(1-Se2)+p1*(1-Sp1)*(1-Sp2)$

- Population 2

	T2+	T2-
T1+	$p2*Se1*Se2 + p2*Sp1*Sp2$	$p2*Se1*(1-Se2)+p2*Sp1*(1-Sp2)$
T1-	$p2*(1-Se1)*Se2+p2*(1-Sp1)*Sp2$	$p2*(1-Se1)*(1-Se2)+p2*(1-Sp1)*(1-Sp2)$

- 6 parameters and 6 equations » Identifiable model!



## Hui-Walter models with multiple populations

- Basically an extension of the single-population model
- Works best with multiple populations each with differing prevalence
- Including an unexposed population works well BUT be wary of assumptions regarding constant sensitivity/specificity across populations with very different types of infections

## Independent intercepts for populations

```
model{  
  for(p in 1:Populations){  
    Tally[1:4, p] ~ dmulti(prob[1:4, p], TotalTests[p])  
    # Test1- Test2- Pop1  
    prob[1, p] <- (prev[p] * ((1-se[1])*(1-se[2]))) + ((1-p  
    ## etc ##  
    prev[p] ~ dbeta(1, 1)  
  }  
  se[1] ~ dbeta(HPSe[1,1], HPSe[1,2])T(1-sp[1], )  
  sp[1] ~ dbeta(HPSp[1,1], HPSp[1,2])  
  se[2] ~ dbeta(HPSe[2,1], HPSe[2,2])T(1-sp[2], )  
  sp[2] ~ dbeta(HPSp[2,1], HPSp[2,2])  
  #data# Tally, TotalTests, Populations, HPSe, HPSp  
  #monitor# prev, prob, se, sp  
  #inits# prev, se, sp
```

We would usually start with individual-level data in a dataframe:

```
se1 <- 0.9; se2 <- 0.8; sp1 <- 0.95; sp2 <- 0.99
prevalences <- c(0.1, 0.5, 0.9)
N <- 100
simdata <- data.frame(Population = sample(seq_along(prevalences), N, replace = TRUE))
simdata$probability <- prevalences[simdata$Population]
simdata$truestatus <- rbinom(N, 1, simdata$probability)
simdata$Test1 <- rbinom(N, 1, (simdata$truestatus * se1) + (1 - simdata$truestatus * se1))
simdata$Test2 <- rbinom(N, 1, (simdata$truestatus * se2) + (1 - simdata$truestatus * se2))
```

```
head(simdata)
```

##	Population	probability	truestatus	Test1	Test2
## 1	3	0.9	1	1	0
## 2	2	0.5	0	0	0
## 3	3	0.9	0	0	0
## 4	2	0.5	1	1	1
## 5	2	0.5	1	1	1
## 6	3	0.9	1	0	0

[Except that probability and truestatus would not normally be known!]

The model code and data format for an arbitrary number of populations (and tests) can be determined automatically

There is a function (provided in the GitHub repo) that can do this for us:

```
simdata$Population <- factor(simdata$Population, levels=source("autohuiwalter.R"))
auto_huiwalter(simdata[,c('Population', 'Test1', 'Test2')],
               outfile='autohw.bug')
```

This generates self-contained model/data/initial values etc (ignore covse and covsp for now):

```
model{

  ## Observation layer:

  # Complete observations (N=100):
  for(p in 1:Populations){
    Tally_RR[1:4,p] ~ dmulti(prob_RR[1:4,p], N_RR[p])

    prob_RR[1:4,p] <- se_prob[1:4,p] + sp_prob[1:4,p]
  }

  ## Observation probabilities:

  for(p in 1:Populations){
```

And can be run directly from R:

```
results <- run.jags('autohw.bug')  
results
```

	Lower95	Median	Upper95	SSeff	psrf
se[1]	0.829	0.944	1.000	5480	1.001
se[2]	0.533	0.688	0.833	9979	1.000
sp[1]	0.889	0.962	1.000	5925	1.000
sp[2]	0.910	0.970	1.000	5883	1.000
prev[1]	0.000	0.070	0.175	7040	1.000
prev[2]	0.222	0.382	0.553	10413	1.000
prev[3]	0.627	0.796	0.936	9387	1.000
covse12	0.000	0.000	0.000	NA	NA
covsp12	0.000	0.000	0.000	NA	NA

- Modifying priors must still be done directly in the model file
- The model needs to be re-generated if the data changes
- But remember that your modified priors will be reset
- There must be a single column for the population (as a factor), and all of the other columns (either factor, logical or numeric) are interpreted as being test results
- The function will soon be included in the runjags package